

Ownership and Borrowing in Rust

Michael Birch

University of Waterloo, May 2023

GitHub Repository

- ❖ Code examples found here



Why Rust?

- ❖ Garbage collection creates overhead
 - ❖ E.g., Java
- ❖ Manual memory management can cause bugs
 - ❖ Memory leaks
 - ❖ Invalid memory access (seg fault)
- ❖ Rust's solution: compile-time “borrow checker”
 - ❖ Bonus: also prevents data races



Drop

- ❖ Prevents memory leaks
- ❖ Enables RAII (Resource Allocation is Initialization)

Drop

```
#[test]
▶ Run Test | Debug
fn test_drop() {
    let _alice: CustomDrop = CustomDrop { name: "Alice" };
    let x: u32 = {
        let _bob: CustomDrop = CustomDrop { name: "Bob" };
        42
    };
    println!("{}");
}

#[derive(Debug)]
2 implementations
pub struct CustomDrop {
    pub name: &'static str,
}

impl Drop for CustomDrop {
    fn drop(&mut self) {
        println!("Dropping {}", self.name);
    }
}
```

```
cargo test test_drop -- --nocapture

running 1 test
Dropping Bob
42
Dropping Alice
test drop::test_drop ... ok
```

Move

- ❖ Variables can be **moved** into another scope or data type

```
#[test]
▶ Run Test | Debug
fn test_move() {
    let _container: Owner<CustomDrop> = {
        let alice: CustomDrop = CustomDrop { name: "Alice" };
        Owner { inner: alice }
    };
    println!("End of scope");
}
```

```
cargo test test_move -- --nocapture
running 1 test
End of scope
Dropping owner container
Dropping Alice
test ownership::test_move ... ok
```

Move

- ❖ Variables can be **moved** into another scope or data type

```
#[test]
▶ Run Test | Debug
fn test_function_move() {
    fn move_var<T: std::fmt::Debug>(value: T) {
        println!("Taken ownership of {value:?}");
    }

    let alice: CustomDrop = CustomDrop { name: "Alice" };
    move_var(alice);
    println!("End of scope");
}
```

```
cargo test test_function_move -- --nocapture
Taken ownership of CustomDrop { name: "Alice" }
Dropping Alice
End of scope
test ownership::test_function_move ... ok
```

Move

- ❖ Variables can be **moved** into another scope or data type
- ❖ A variable must have a single owner to prevent invalid memory access (seg fault)

```
fn test_double_move() {  
    let alice: CustomDrop = crate::drop::CustomDrop { name: "Alice" };  
    let _owner1: Owner<CustomDrop> = crate::ownership::Owner {inner: ...};  
    let _owner2: Owner<CustomDrop> = crate::ownership::Owner {inner: alice};  
}
```

```
Compiling borrowing-and-ownership v0.1.0 (/mnt/c/Users/mbirc/Documents/rust/borrowing_and_ownership)  
error[E0382]: use of moved value: `alice`  
--> src/ownership.rs:51:33  
  
49 |     let alice = crate::drop::CustomDrop { name: "Alice" };  
50 |         ----- move occurs because `alice` has type `CustomDrop`, which does not implement the `Copy` trait  
51 |     let _owner1 = Owner {inner: alice};  
                  ----- value moved here  
51 |     let _owner2 = Owner {inner: alice};  
                  ^^^^^^ value used here after move
```

Copy

- ❖ Tells the compiler to copy a value if it is moved multiple times
- ❖ Automatically implemented for primitive types like u32

```
#[test]
▶ Run Test | Debug
fn test_copy() {
    let alice: Copyable = Copyable { name: "Alice" };
    let _owner1: Owner<Copyable> = crate::ownership::Owner { inner: alice };
    let _owner2: Owner<Copyable> = crate::ownership::Owner { inner: alice };
    println!("End of scope");
}
```

Borrowing

- ❖ Borrowing allows accessing a value without moving or copying it

#[test] ❖ Borrowed values are called “references”

▶ Run Test | Debug

```
fn test_borrow() {
    fn borrow_to_print<T: std::fmt::Debug>(value: &T) {
        println!("Borrowed {value:?}");
    }
}
```

```
let alice: CustomDrop = crate::drop::CustomDrop { name: "Alice" };
borrow_to_print(&alice);
borrow_to_print(&alice);
println!("End of scope");
}
```

```
cargo test test_borrow -- --nocapture
running 1 test
Borrowed CustomDrop { name: "Alice" }
Borrowed CustomDrop { name: "Alice" }
End of scope
Dropping Alice
test borrow::test_borrow ... ok
```

Borrowing

- ❖ Data types can hold references, but require a “lifetime” parameter

```
0 implementations
pub struct Borrower<'a, T> {
    pub value: &'a T,
}

#[test]
▶ Run Test | Debug
fn test_data_type_borrow() {
    let alice: CustomDrop = crate::drop::CustomDrop { name: "Alice" };
    let _x: u32 = {
        let _borrowed1: Borrower<CustomDrop> = Borrower { value: &alice };
        let _borrowed2: Borrower<CustomDrop> = Borrower { value: &alice };
        42
    };
    println!("End of scope");
}
```

Borrowing

- ❖ Data types can hold references, but require a “lifetime” parameter
- ❖ Lifetimes are how the compiler keeps track of when values are dropped to prevent seg faults

```
fn test_value_does_not_live_long_enough() {
    let _borrow = {
        let alice = crate::drop::CustomDrop { name: "Alice" };
        Borrower { value: &alice }
    };
}

error[E0597]: `alice` does not live long enough
--> src/borrow.rs:52:27
|
50 |     let _borrow = {
51 |         ----- borrow later stored here
52 |         let alice = crate::drop::CustomDrop { name: "Alice" };
53 |             ----- binding `alice` declared here
|         Borrower { value: &alice }
|                         ^^^^^^ borrowed value does not live long enough
|         };
|         - `alice` dropped here while still borrowed
```

Cloning

- ❖ **Clone** is similar to **Copy**, but the compiler does not automatically make the copy
- ❖ **Clone** operations are supposed to be much more expensive than **Copy**
 - ❖ (e.g. require heap allocation)
- ❖ **Clone** operates on a reference

```
pub trait Clone {  
    fn clone(&self) -> Self;  
}
```

Cloning

- ❖ **Clone** is similar to **Copy**, but the compiler does not automatically make the copy
- ❖ **Clone** operations are supposed to be much more expensive than **Copy**
 - ❖ (e.g. require heap allocation)
- ❖ **Clone** operates on a reference

```
#[test]
▶ Run Test | Debug
fn test_clone() {
    let alice: Clonable = Clonable {
        name: "Alice".into(),
    };
    let _owner1: Owner<Clonable> = crate::ownership::Owner {
        inner: alice.clone(),
    };
    let _owner2: Owner<Clonable> = crate::ownership::Owner { inner: alice };
    println!("End of scope");
}
```

Reference Counting

- ❖ Rc (and Arc) is a special type in the standard library to “cheat” the ownership rule
- ❖ The object is owned by an allocation on the heap; it is dropped when the last reference object is dropped.
- ❖ Cloning references does not duplicate the underlying object

Reference Counting

```
#[test]
▶ Run Test | Debug
fn test_reference_counted() {
    use std::rc::Rc;

    let alice: CustomDrop = crate::drop::CustomDrop {name: "Alice"};
    let rc1: Rc<CustomDrop> = Rc::new(alice);
    let rc2: Rc<CustomDrop> = Rc::clone(self: &rc1);
    let rc3: Rc<CustomDrop> = Rc::clone(self: &rc1);
    drop(rc1);
    println!("Dropped one reference");
    drop(rc2);
    println!("Dropped second reference");
    drop(rc3);
    println!("Dropped final reference");
}
```

```
cargo test test_reference_counted -- --nocapture
running 1 test
Dropped one reference
Dropped second reference
Dropping Alice
Dropped final reference
test reference_counted::test_reference_counted ... ok
```

Mutable Borrowing

- ❖ Mutable references are distinct from immutable ones
- ❖ Any number of immutable references are allowed; only one mutable is allowed

```
#[test] ❖ This prevents data races
```

```
▶ Run Test | Debug
```

```
fn test_mut_borrow() {
    fn transmute_to_bob(value: &mut CustomDrop) {
        value.name = "Bob";
    }
}
```

```
let mut alice: CustomDrop = CustomDrop {name: "Alice"};
transmute_to_bob(&mut alice);
println!("End of scope");
}
```

```
cargo test test_mut_borrow -- --nocapture
running 1 test
End of scope
Dropping Bob
test borrow::test_mut_borrow ... ok
```

Mutable Borrowing

- ❖ Mutable references are distinct from immutable ones
- ❖ Any number of immutable references are allowed; only one mutable is allowed

```
# [test] ❖ This prevents data races
▶ Run Test | Debug
fn test_mut_borrow() {
    fn transmute_to_bob(value: &mut CustomDrop) {
        value.name = "Bob";
    }

    let mut alice: CustomDrop = CustomDrop {name: "Alice"};
    let reference: &CustomDrop = &alice;
    transmute_to_bob(&mut alice);      error[E0502]: cannot borrow `alice` as mutable because it is also borrowed as immutable
    println!("{}{reference:?}");        --> src/borrow.rs:61:22
    println!("End of scope");
}

60 |     let reference = &alice;
61 |             ----- immutable borrow occurs here
62 |     transmute_to_bob(&mut alice);
63 |             ^^^^^^^^^^ mutable borrow occurs here
64 |     println!("{}{reference:?}");
65 |             ----- immutable borrow later used here
```

Interior mutability

- ❖ Allows “cheating” the Rust type system by mutating an object via an immutable reference
- ❖ Sometimes necessary in complex concurrent applications

```
fn test_data_race() {
    use std::{thread, time::Duration};

    // Define a cell with value equal to 40
    let cell: SimpleCell = SimpleCell::new(40);
    let borrow: &SimpleCell = &cell;
    // Define a function which will increment the value of the cell by 1.
    let add_one: impl Fn() = || {
        let mut x: u32 = borrow.get();
        x += 1;
        thread::sleep(dur: Duration::from_millis(10));
        borrow.set(new_value: x);
    };
    std::thread::scope(|s: &Scope| {
        // Create two threads, both incrementing the cell by 1.
        let thread1: ScopedJoinHandle<()> = s.spawn(add_one);
        let thread2: ScopedJoinHandle<()> = s.spawn(add_one);

        thread1.join().unwrap();
        thread2.join().unwrap();
    });
    assert_eq!(cell.get(), 42);
}
```

- ❖ Using unsafe code we can write an example of a data race

```
thread 'interior_mut::test_data_race' panicked at
'assertion failed: `(left == right)''  

  left: '41',  

  right: '42'', src/interior_mut.rs:34:5
```

- ❖ Using types in the standard library, interior mutability can be done safely

```
fn test_safe_interior_mutability() {
    use std::{sync::RwLock, thread, time::Duration};

    let cell: RwLock<u32> = RwLock::new(40_u32);
    let borrow: &RwLock<u32> = &cell;
    let add_one: impl Fn() = || {
        let mut x: RwLockWriteGuard<u32> = borrow.write().unwrap();
        thread::sleep(dur: Duration::from_millis(10));
        *x += 1;
    };
    std::thread::scope(|s: &Scope| {
        // Create two threads, both incrementing the cell by 1.
        let thread1: ScopedJoinHandle<()> = s.spawn(add_one);
        let thread2: ScopedJoinHandle<()> = s.spawn(add_one);

        thread1.join().unwrap();
        thread2.join().unwrap();
    });
    assert_eq!(*cell.read().unwrap(), 42);
}
```

```
cargo test test_safe_interior_mutability -- --nocapture
running 1 test
test interior_mut::test_safe_interior_mutability ... ok
```

Exercise: Map-Reduce

```
pub fn map_red<T, U, M, R>(data: &[T], map_fn: M, reduce: R, init: U) -> U
where
    // Don't worry about the Send + Sync everywhere, that's just to tell rust the data
    // can be shared between threads safely.
    T: Send + Sync,
    U: Clone + Send + Sync,
    M: Fn(&T) -> U + Send + Sync,
    R: Fn(U, U) -> U + Send + Sync,
{
    data.iter().map(map_fn).fold(init, f: reduce)
}
```

Solution: Map-Reduce

```
pub fn map_red<T, U, M, R>(data: &[T], map_fn: M, reduce: R, init: U) -> U
where ...
{
    //data.iter().map(map_fn).fold(init, reduce)
    let chunk_size: usize = data.len() / 2;
    let process_chunk: impl Fn(&[T], U) -> U =
        |input: &[T], init: U| -> U { input.iter().map(&map_fn).fold(init, f: &reduce) };
    std::thread::scope(|s: &Scope| {
        let init_copy: U = init.clone();
        let handle_1: ScopedJoinHandle<U> = s.spawn(|| process_chunk(input: &data[0..chunk_size], init_copy));
        let handle_2: ScopedJoinHandle<U> = s.spawn(|| process_chunk(input: &data[chunk_size..], init));
        let result_1: U = handle_1.join().unwrap();
        let result_2: U = handle_2.join().unwrap();
        reduce(result_1, result_2)
    })
}
```

Questions?

Contact

- ❖ @birchmd (Telegram)
- ❖ <https://github.com/birchmd/>
- ❖ <https://www.typedriven.ca/news/>
 - ❖ Occasional blog about Rust / blockchain stuff